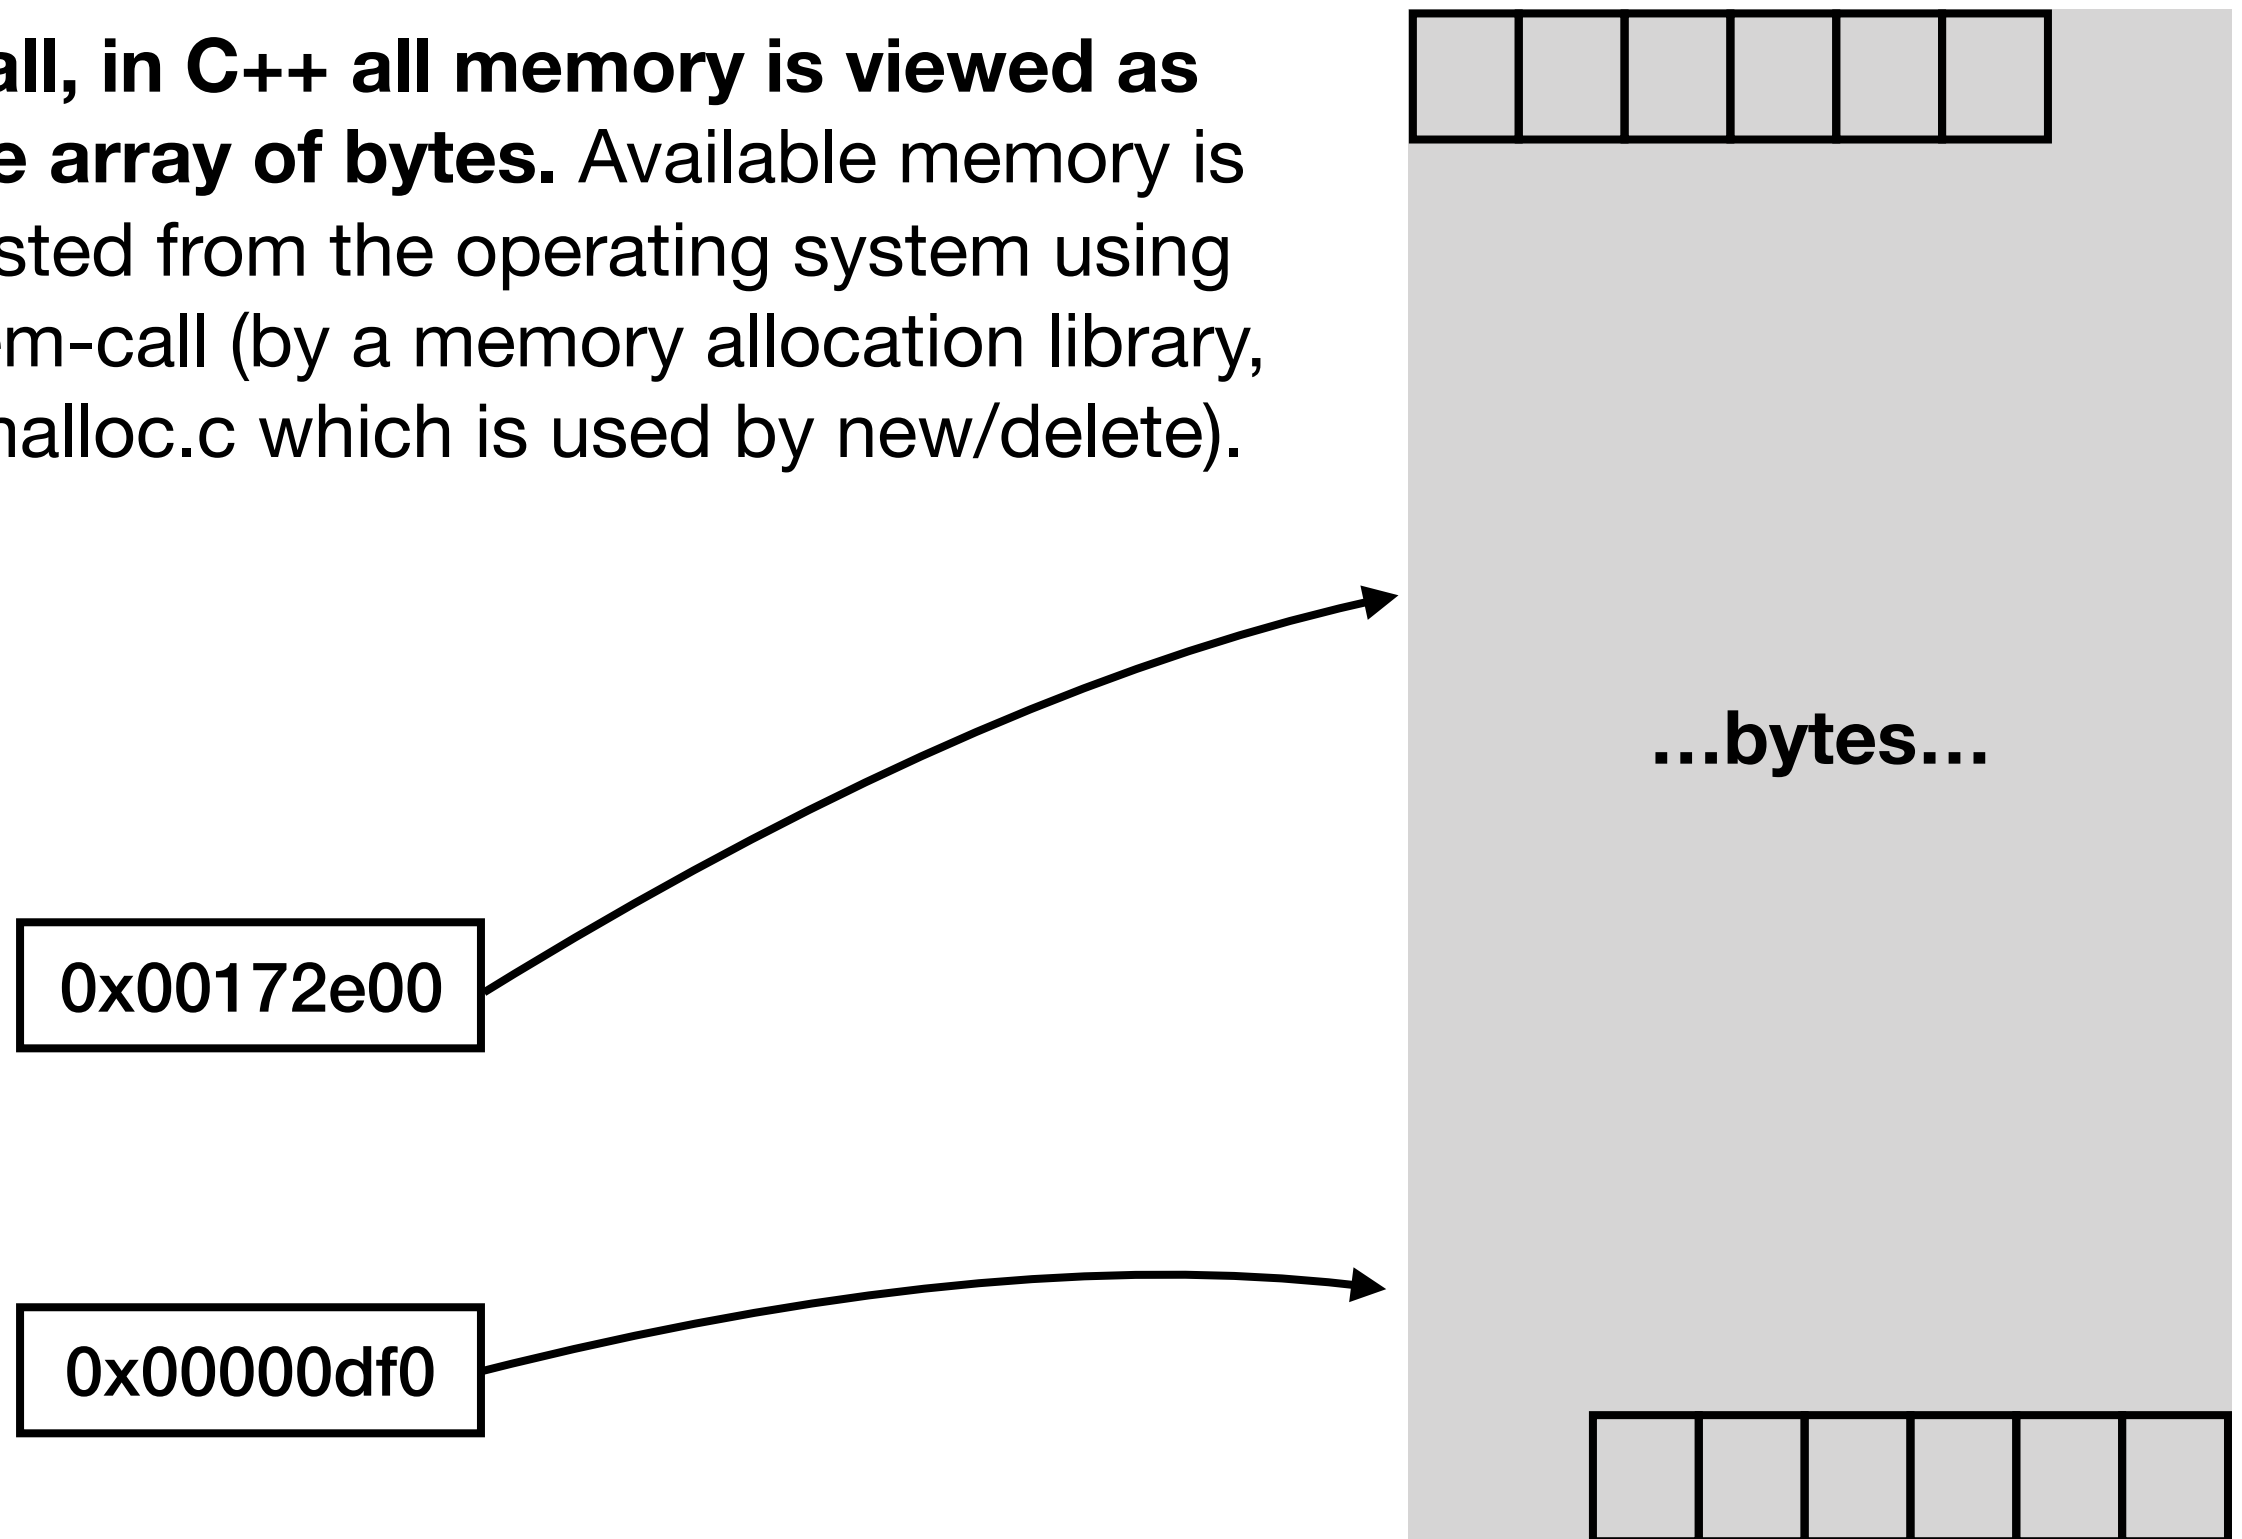


Memory Management

C++ semantics: memory model

Recall, in C++ all memory is viewed as a huge array of bytes. Available memory is requested from the operating system using a system-call (by a memory allocation library, e.g., malloc.c which is used by new/delete).



C++ semantics: memory model

Recall, in C++ all memory is viewed as a huge array of bytes. Available memory is requested from the operating system using a system-call (by a memory allocation library, e.g., malloc.c which is used by new/delete).

0x00172e00


0x00000df0

The stack starts growing down!
The C++ runtime reserves a portion of memory (that is extended dynamically upon a page fault).



The heap starts below the stack in memory and grows up, page by page.

C++ semantics: memory model



Recall, in C++ pointers are (virtual) memory addresses and refer to the start of a buffer.

Exactly how many bytes are being used by this pointer, after that location, is determined by how the C++ program uses that pointer! **(E.g., C-strings are null-value terminated.)**

This is not statically checked, leading to buffer overflow.

The stack starts growing down!
The C++ runtime reserves a portion of memory (that is extended dynamically upon a page fault).

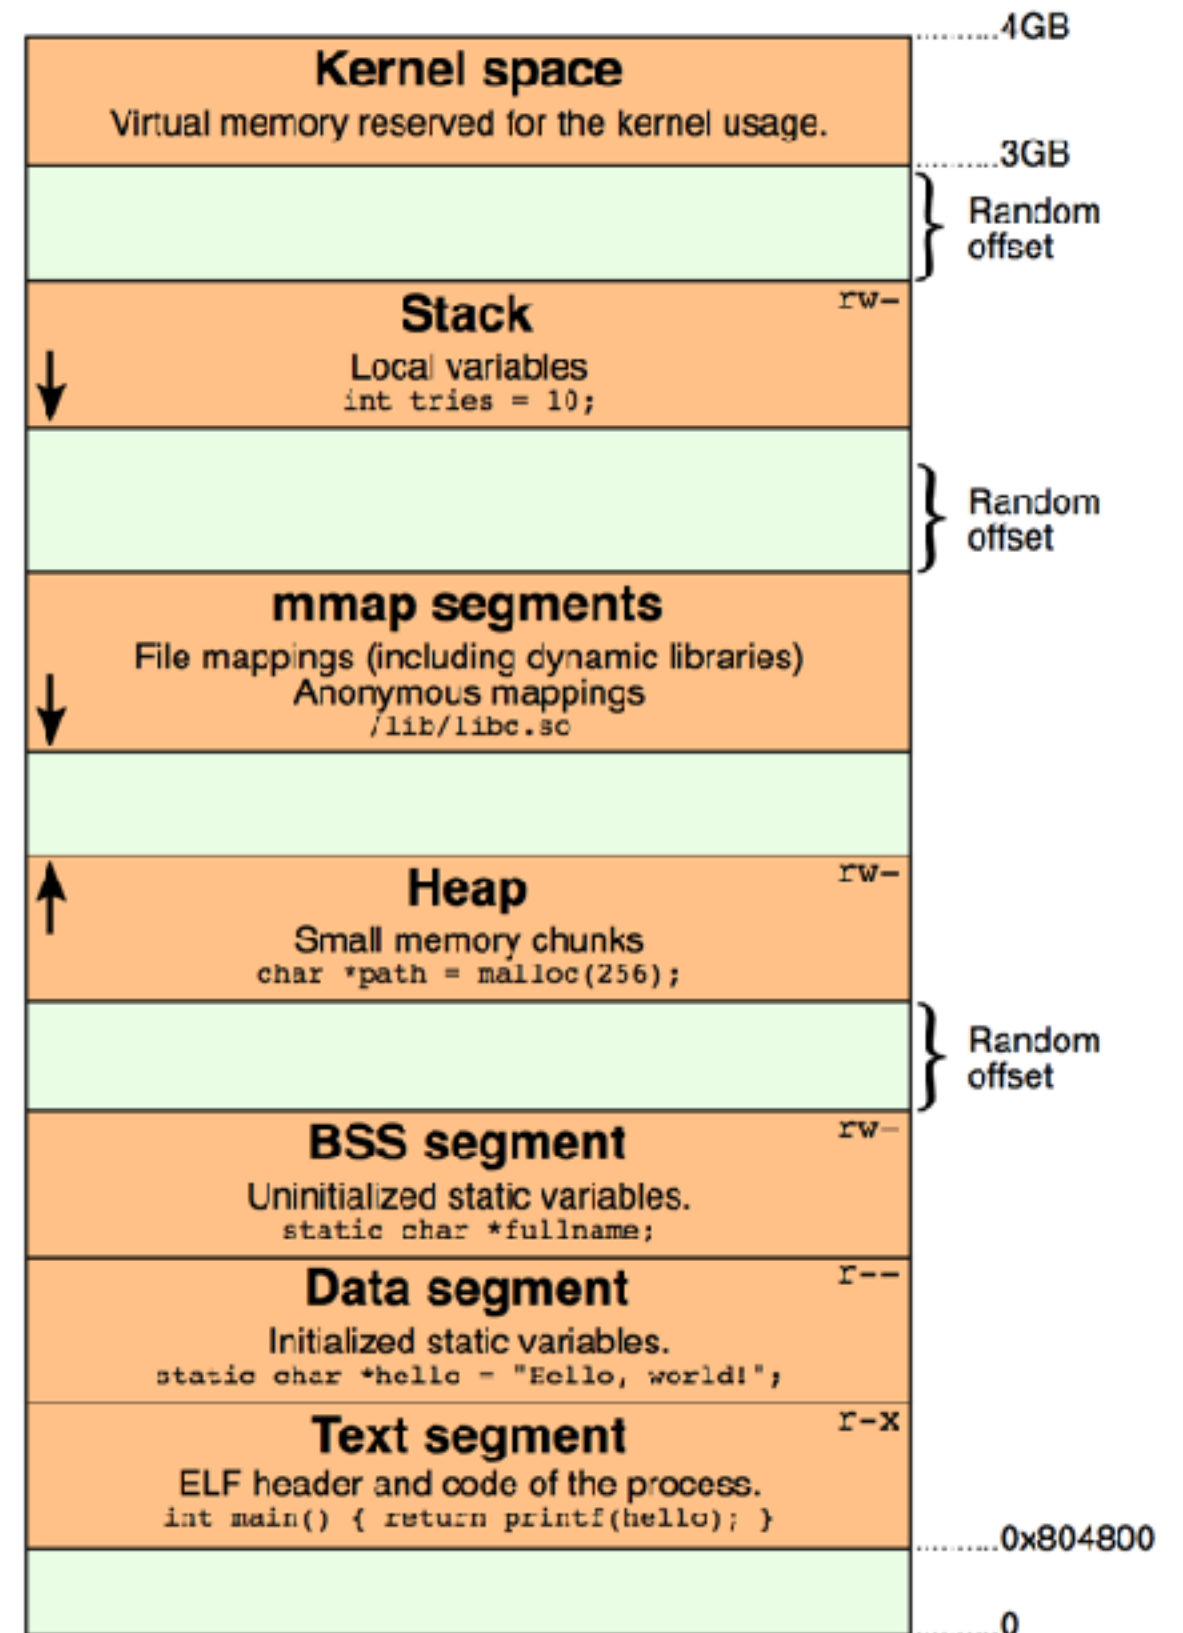


The heap starts below the stack in memory and grows up, page by page.

C++ semantics: memory model

The virtual memory for your C++ binary is organized like so:

Note: The stack grows down. The heap grows up (and is managed by a memory allocator such as malloc in libc).



C++ semantics: pointers and references

**Prefix * operation
turns a pointer into
a reference! *x references
the value at address x.**

```
int* x = f(); // x is a pointer to an int
int y = *x;   // *x dereferences the ptr
```

```
int x = f(); // x is an int
int* y = &x; // &x takes address of x
```

**Prefix & operation
turns a reference into
a pointer! &x is the
address of the value
referenced by x.**

C++ semantics: field access, . and ->

```
A& a = f(); // a is a reference to an object  
int y = a.y; // a.y accesses field y of a
```

The . operation restricts a reference to a specific field; here, a . y turns a reference to a an object into a reference to its y field.

The -> operation dereferences a pointer and accesses a specific field all at once.

```
A* a = f(); // a is a pointer to an object  
int y = a->y; // a->y accesses field y off a
```

C++ semantics: indexing and dereference

**Postfix [..] operation
turns a pointer into
a reference to the
element specified as
the index**

```
int* x = f(); // x is a pointer to an int  
int y = x[0]; // x[0] indexes the pointer
```

```
int* x = f(); // x is a pointer to an int  
int y = *x;    // this is the same as x[0]
```

**If the index is 0, then
this is just the same as
dereferencing the
pointer!**

C++ semantics: indexing and dereference

```
int* x = f();    // x is a pointer to an int
int y = *(x+3); // this is the same as x[3]
```

**If the index is non-0,
then this is just the
same as incrementing
the pointer and then
dereferencing**

**This is the same as
incrementing the raw
address by the
appropriate number
of bytes. The void*
type gives access to
the raw address.**

```
int* x = f();
// x is a pointer to an int
int y = *(int*)((void*)x
               + 3*sizeof(int));
// this is ALSO the same as x[3]
```

C++ semantics: Try an example!

```
int arr[8] = {0,5,1,2,3,4,5,9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << arr[1] << std::endl;  
// Which value is printed out?
```

C++ semantics: Try an example!

```
int arr[8] = {0, 5, 1, 2, 3, 4, 5, 9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << arr[1] << std::endl;  
// Which value is printed out?
```

Answer: 5

C++ semantics: Try an example!

```
int arr[8] = {0,5,1,2,3,4,5,9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << &arr << std::endl;  
// Which value is printed out?
```

C++ semantics: Try an example!

```
int arr[8] = {0,5,1,2,3,4,5,9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << &arr << std::endl;  
// Which value is printed out?
```

Answer: 0xff443120 **<— ptr to var x**
in other words, $(&arr) == 0$**

C++ semantics: Try an example!

```
int arr[8] = {0,5,1,2,3,4,5,9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << (&arr[3])+1 << std::endl;  
// Which value is printed out?
```

C++ semantics: Try an example!

```
int arr[8] = {0,5,1,2,3,4,5,9};  
int* x = arr;           // Derive a ptr from arr  
std::cout << (&arr[3])+1 << std::endl;  
// Which value is printed out?
```

Answer: 0xecff6604 \leftarrow ptr to elem 3
in other words, $*((\&\text{arr}[3])+1) == 3$

C++ semantics: Try an example!

```
int arr[8] = {0, 5, 1, 2, 3, 4, 5, 9};  
int* x = arr;  
int* y = x+3;  
int z = *y;           // What is z?  
int* a = &(y[z]);  
void* b = (void*)a + sizeof(int);  
int c = *((int*)b - 2); // What is c?
```


C++ semantics: Try an example!

```
int arr[8] = {0, 5, 1, 2, 3, 4, 5, 9};  
int* x = arr;  
int* y = x+3;  
int z = *y; // What is z?  
int* a = &(y[z]);  
void* b = (void*)a + sizeof(int);  
int c = *((int*)b - 2); // What is c?
```

Answer: z == 2 && c == 3

reverse.cpp solution

```
struct linkedlist
{
    int value;
    linkedlist* next;
};
```

```
int main()
{
    linkedlist* node = 0; //root
    int n;
    while (std::cin >> n)
    {
        linkedlist* next = node;
        node = new linkedlist();
        node->value = n;
        node->next = next;
    } //...
```

data layout in memory

value	next
int	linkedlist*


```
linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}
```

Calling Function


How C++ sees the **stack**

.....

How C++ sees the **heap**



```
linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}
```




Calling Function


How C++ sees the **stack**



How C++ sees the **heap**



```
linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}
```




Calling Function

How C++ sees the **stack**



How C++ sees the **heap**

```
linkedlist* node = 0; //root
→ int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}
```




Calling Function

How C++ sees the **stack**



How C++ sees the **heap**

```
    linkedlist* node = 0; //root
    int n;
    while (std::cin >> n)
    {
        linkedlist* next = node;
        node = new linkedlist();
        node->value = n;
        node->next = next;
    }
```



Calling Function

How C++ sees the **stack**





How C++ sees the **heap**

```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```



How C++ sees the **stack**

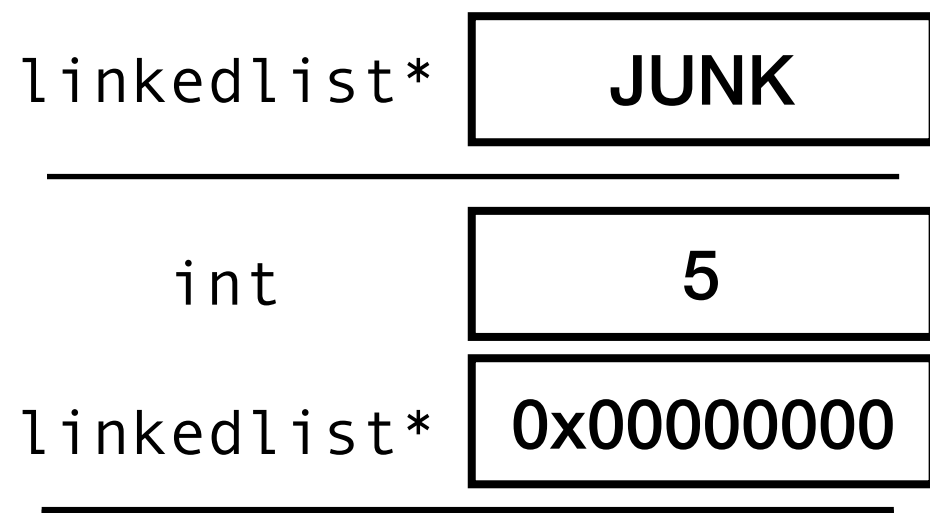


How C++ sees the **heap**


```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
    {
        linkedlist* next = node;
        node = new linkedlist();
        node->value = n;
        node->next = next;
    }

```



Calling Function

How C++ sees the **stack**

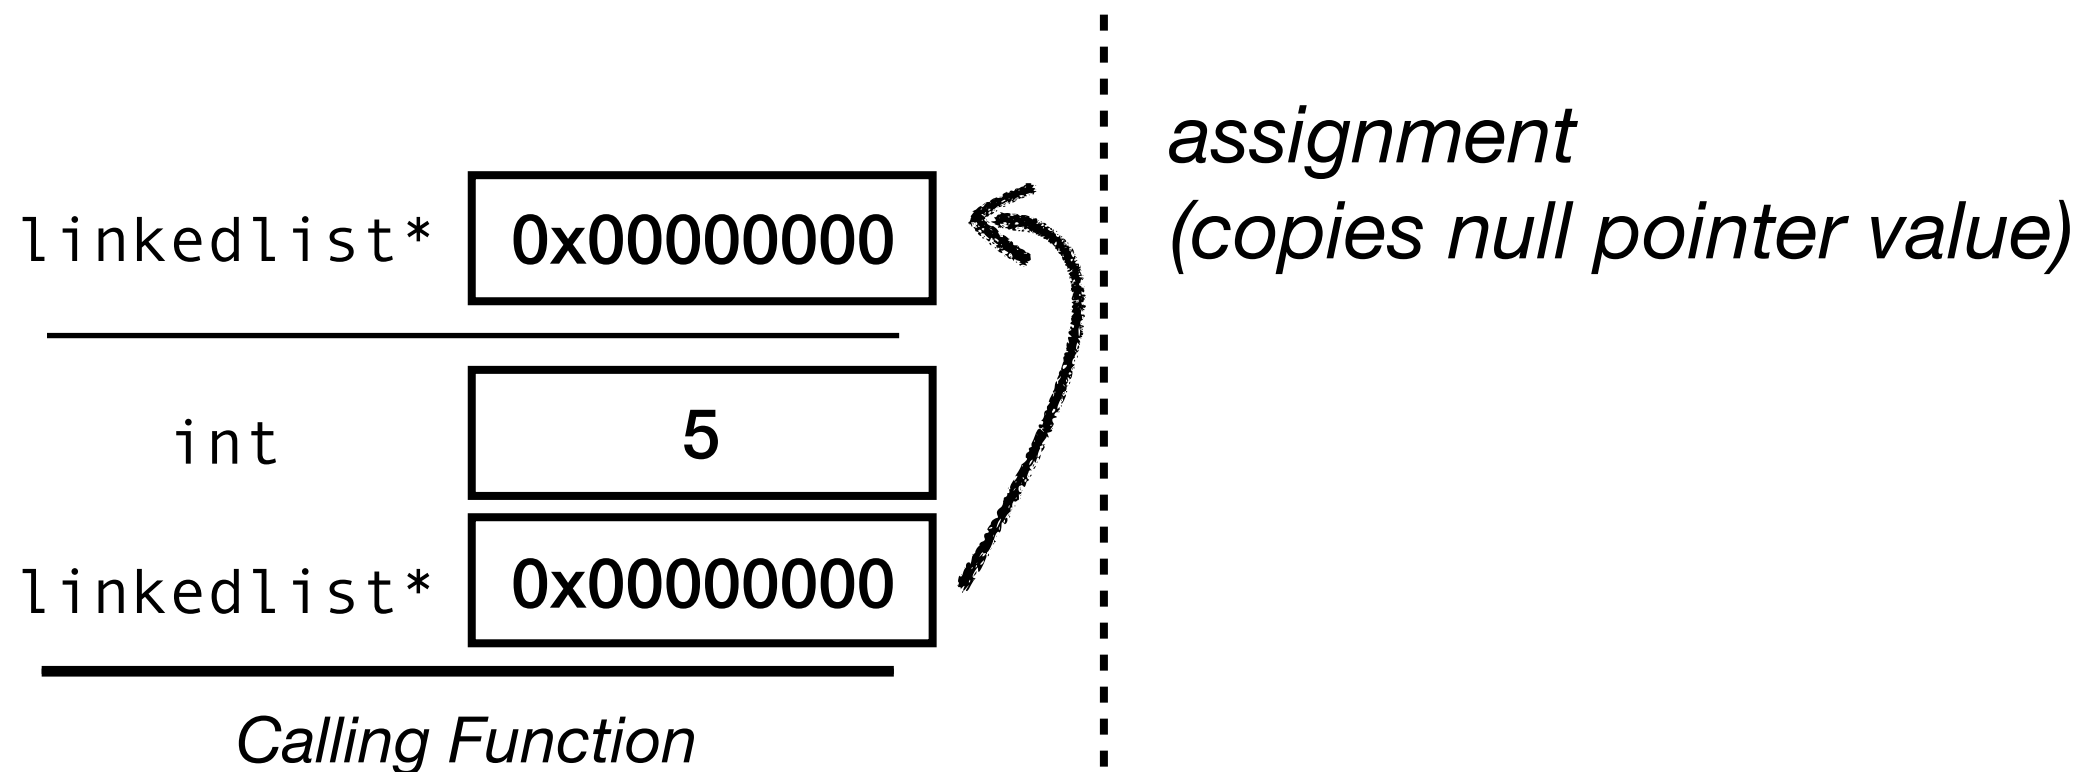



How C++ sees the **heap**

```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}



```



```


linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```



Calling Function

How C++ sees the **stack**

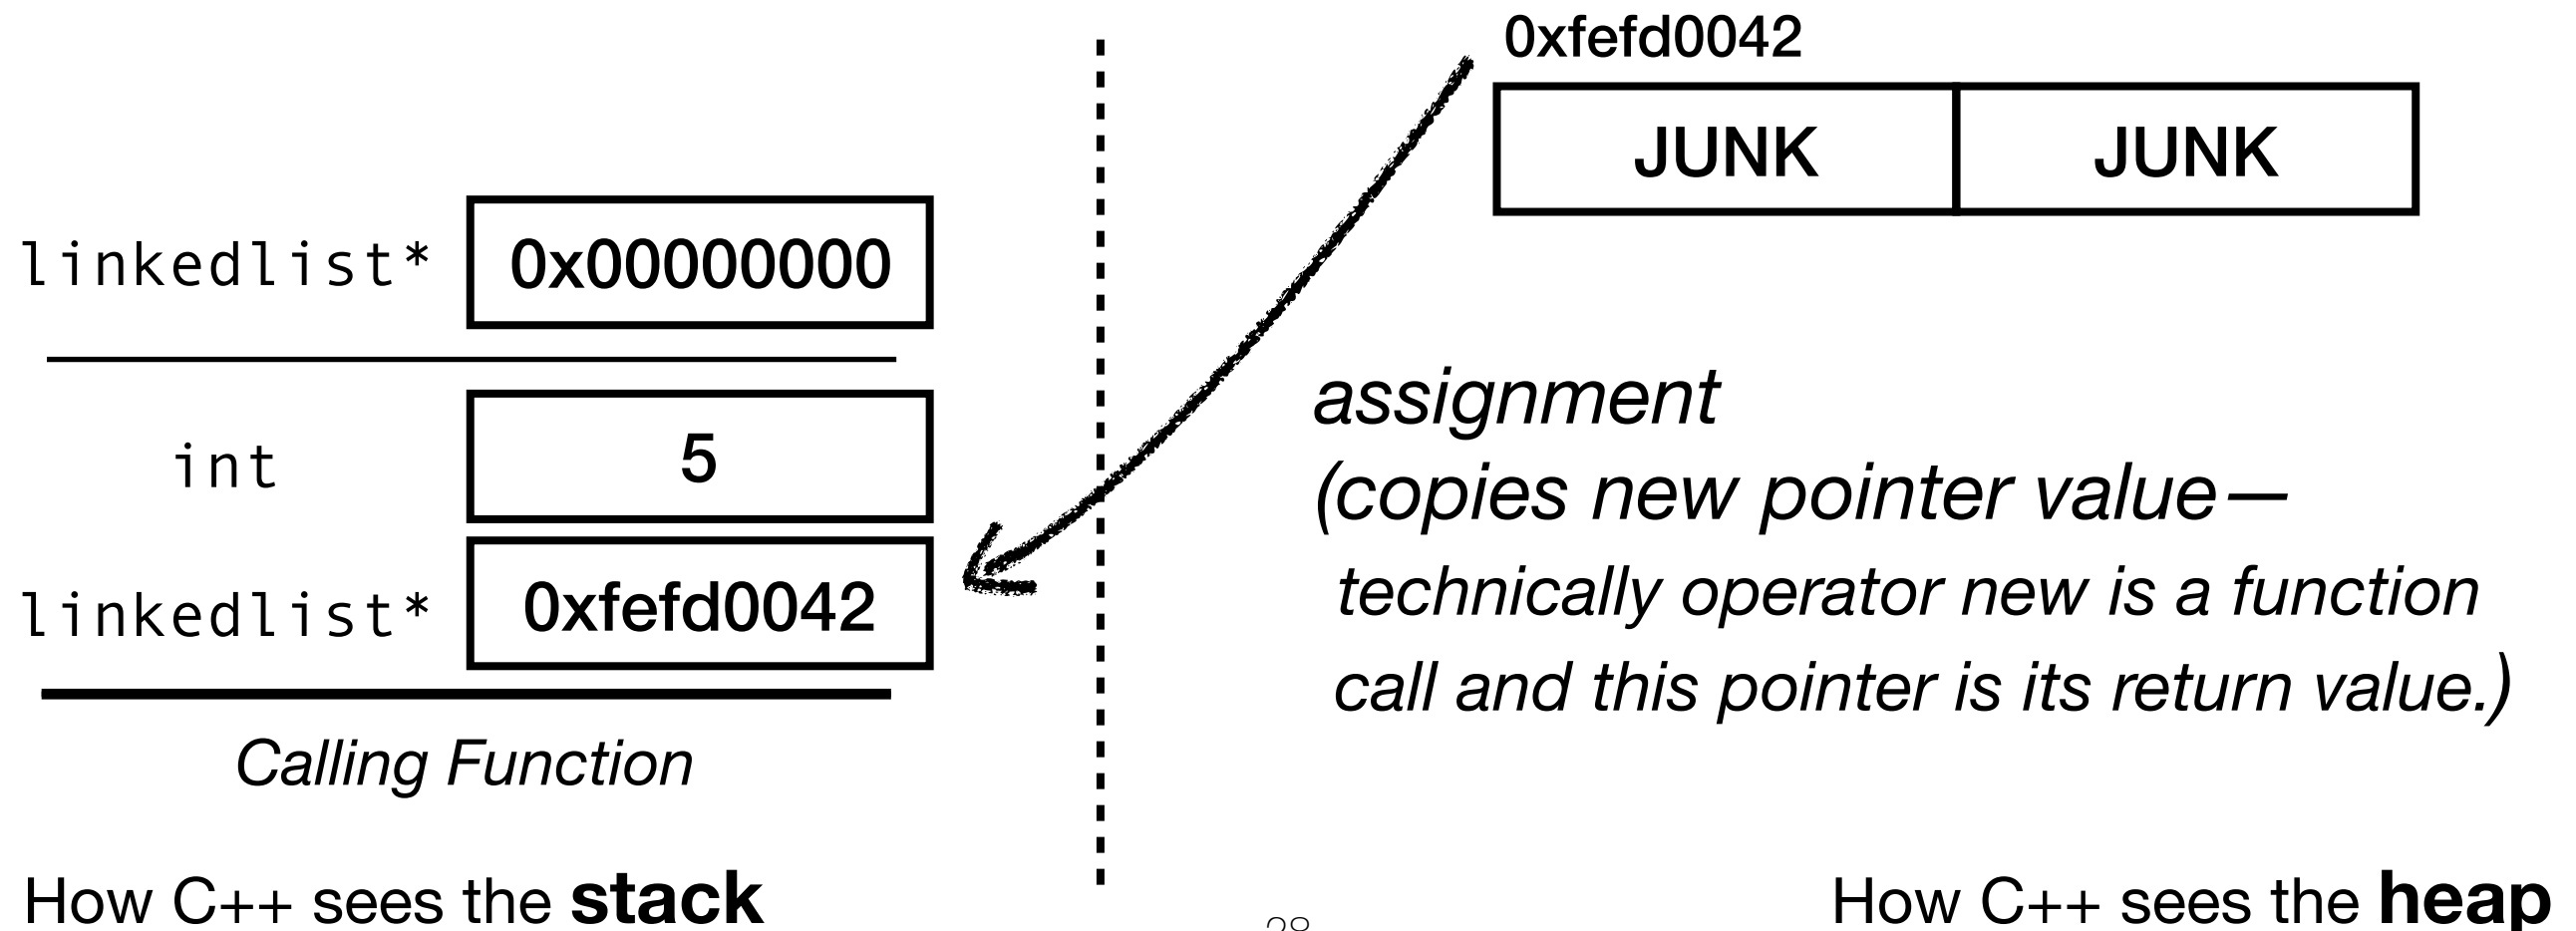



How C++ sees the **heap**

```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}



```



```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

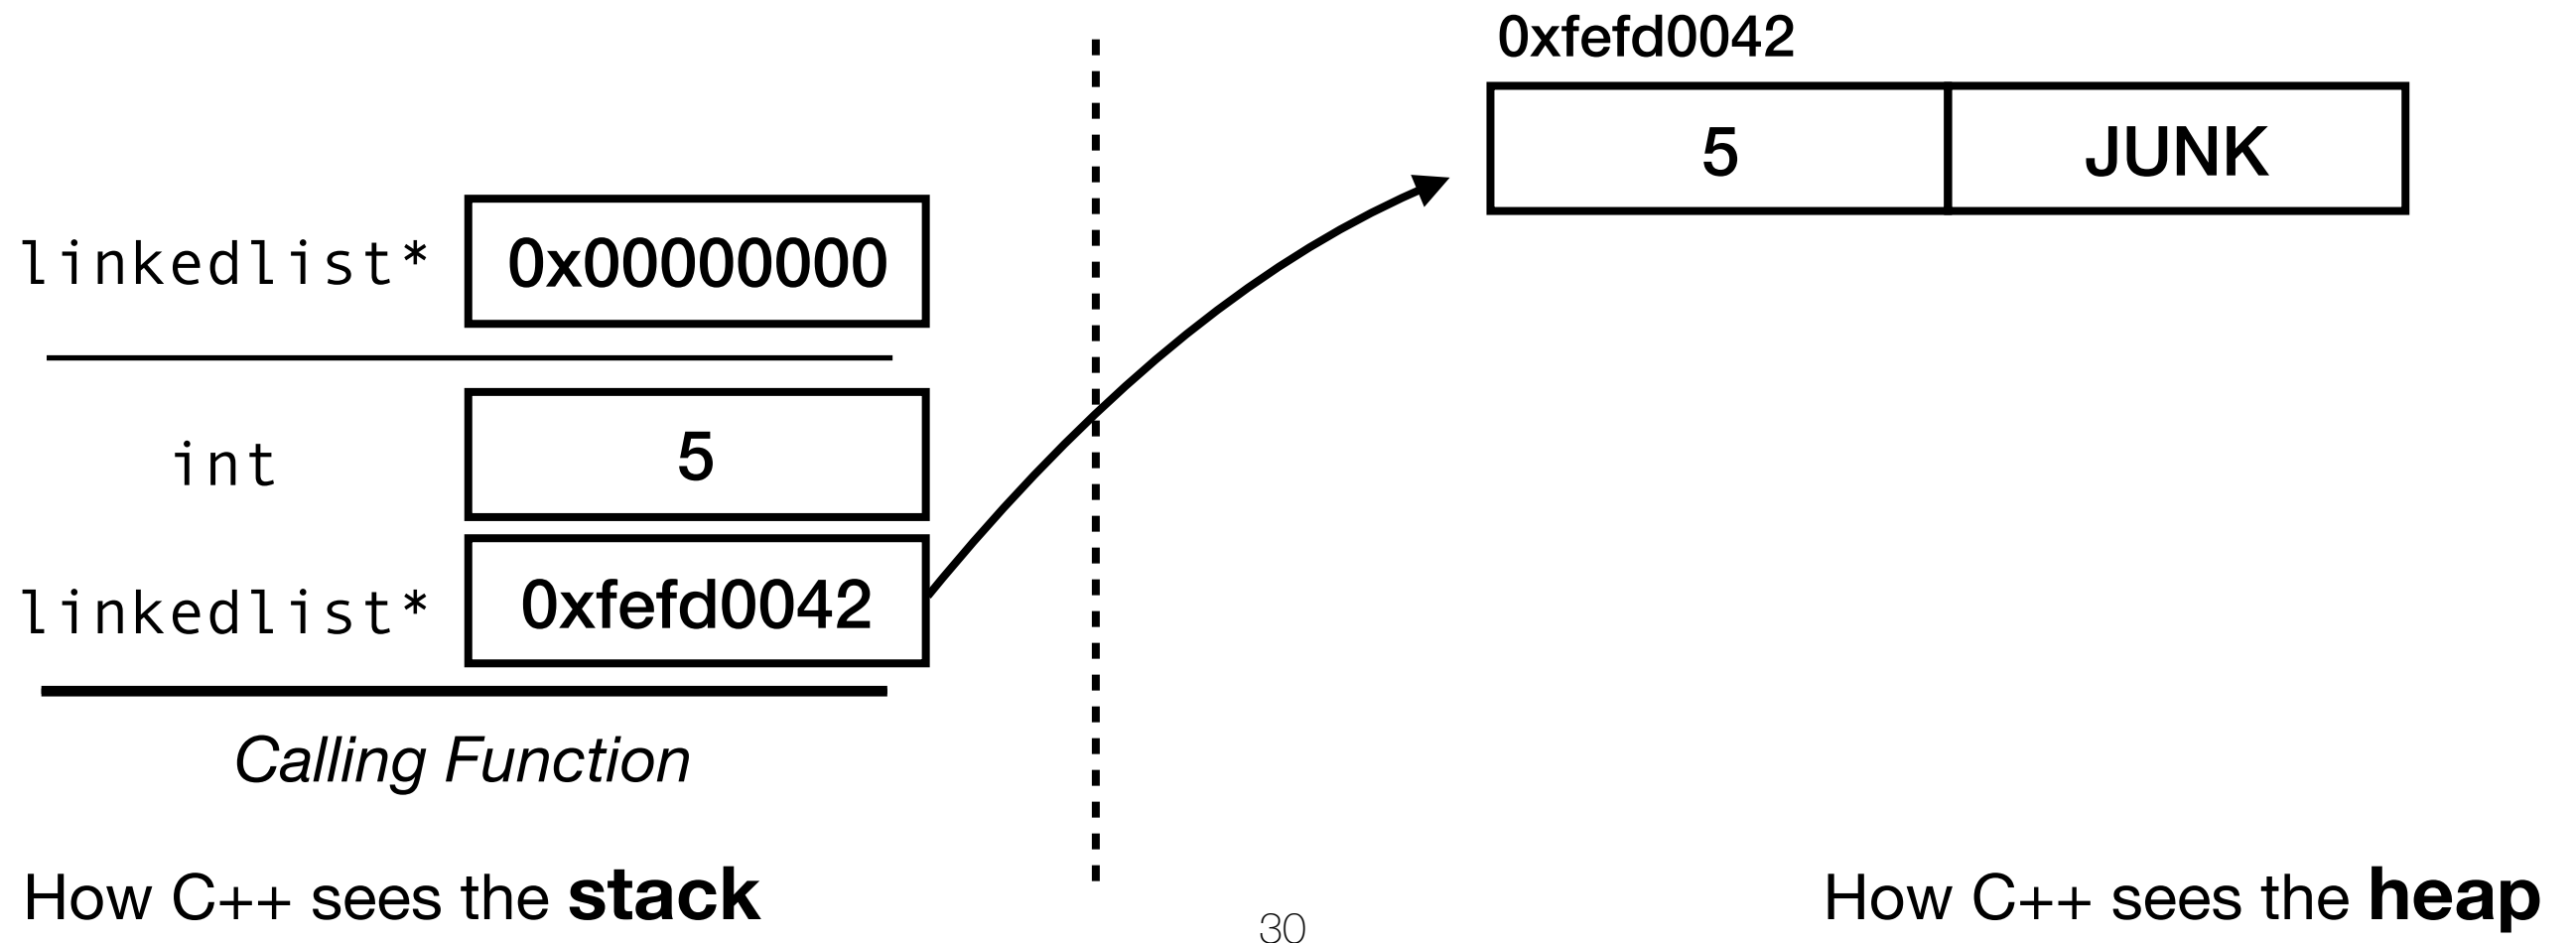

```



```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

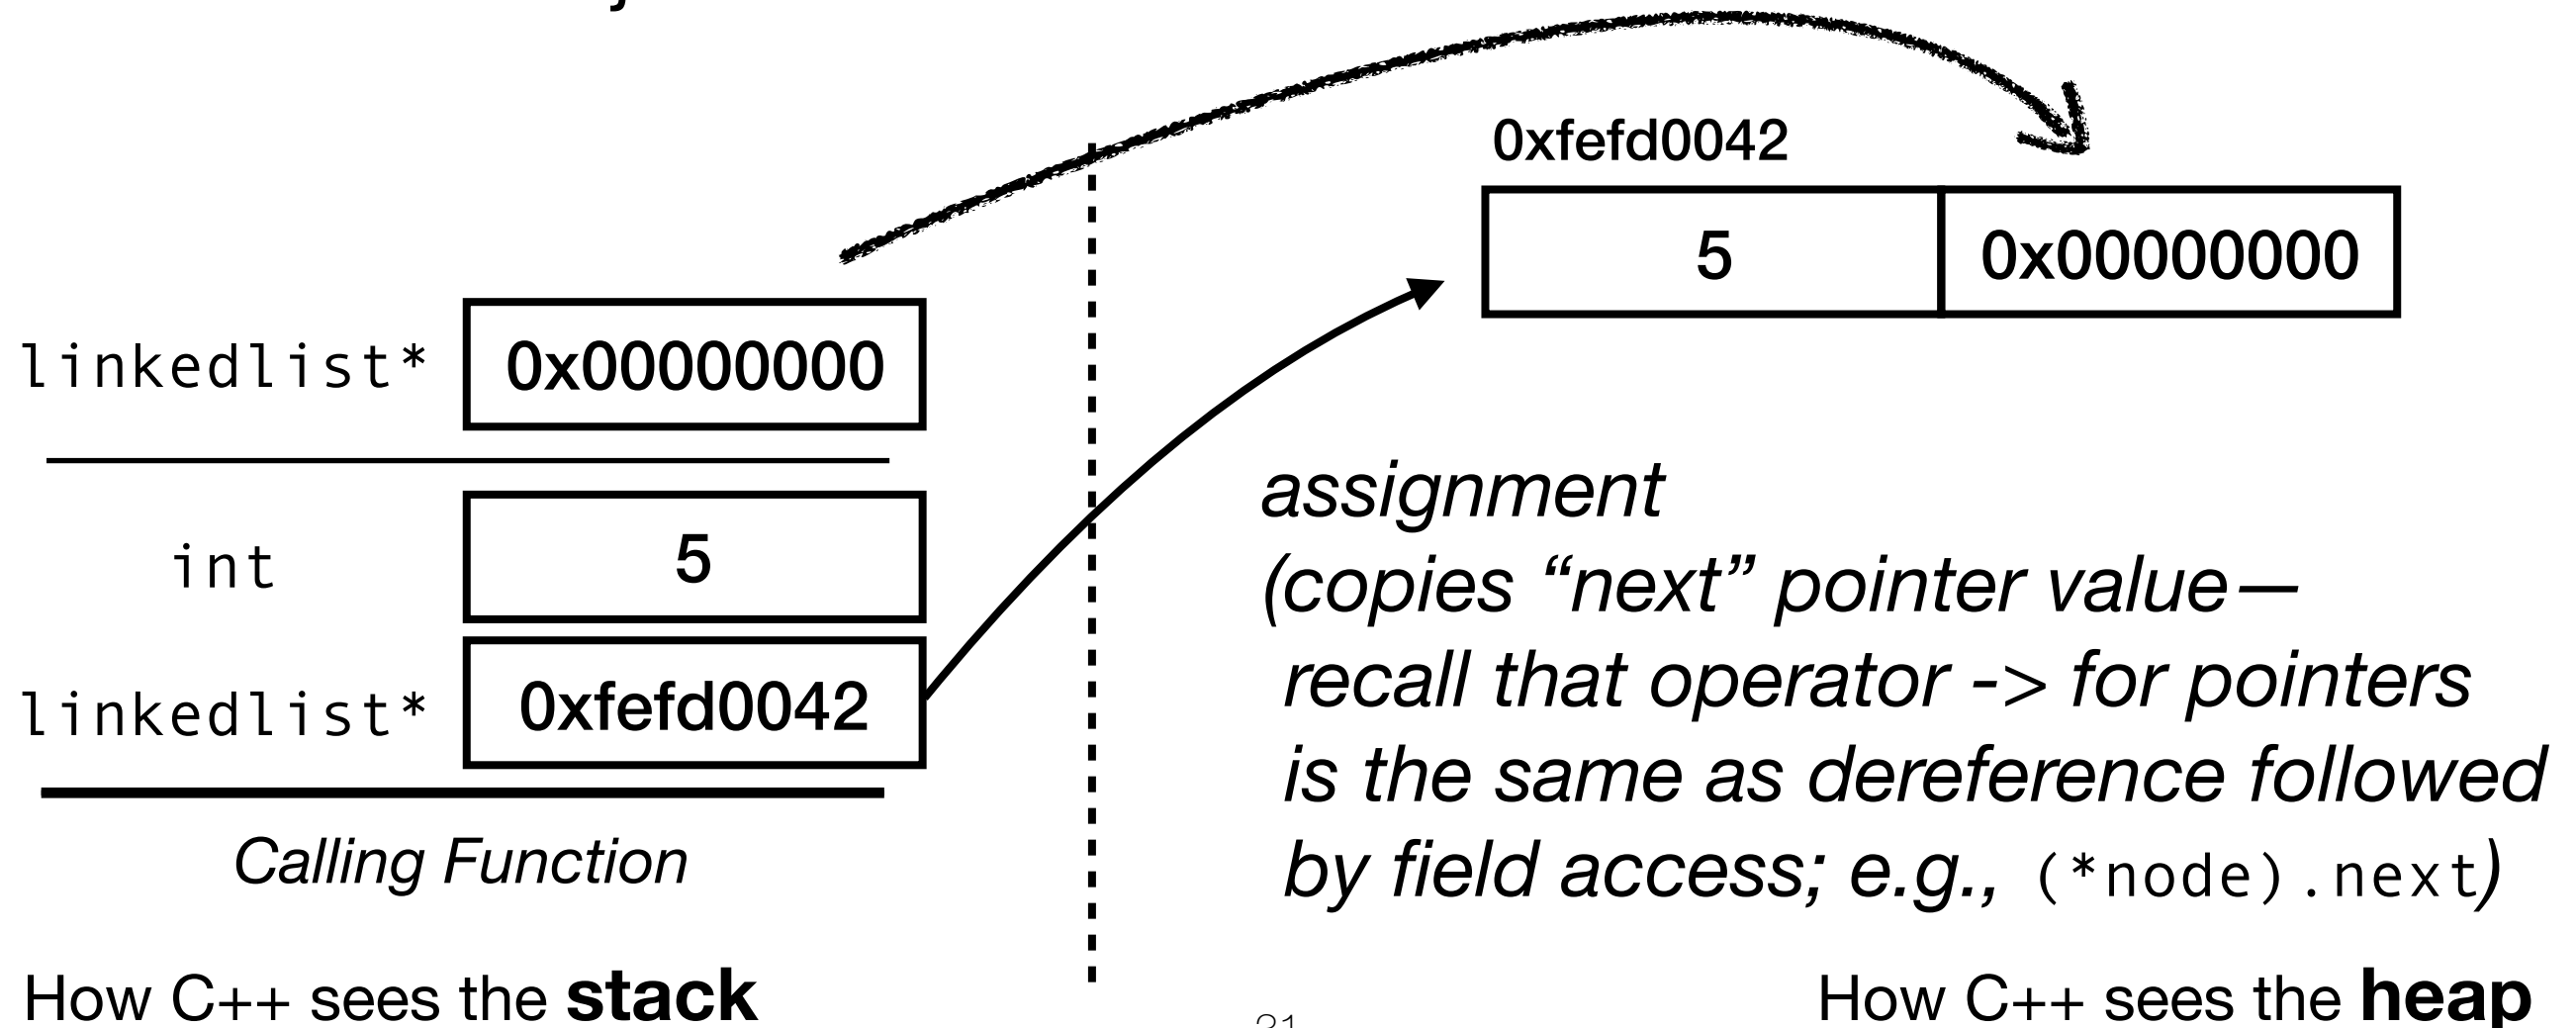

```



```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

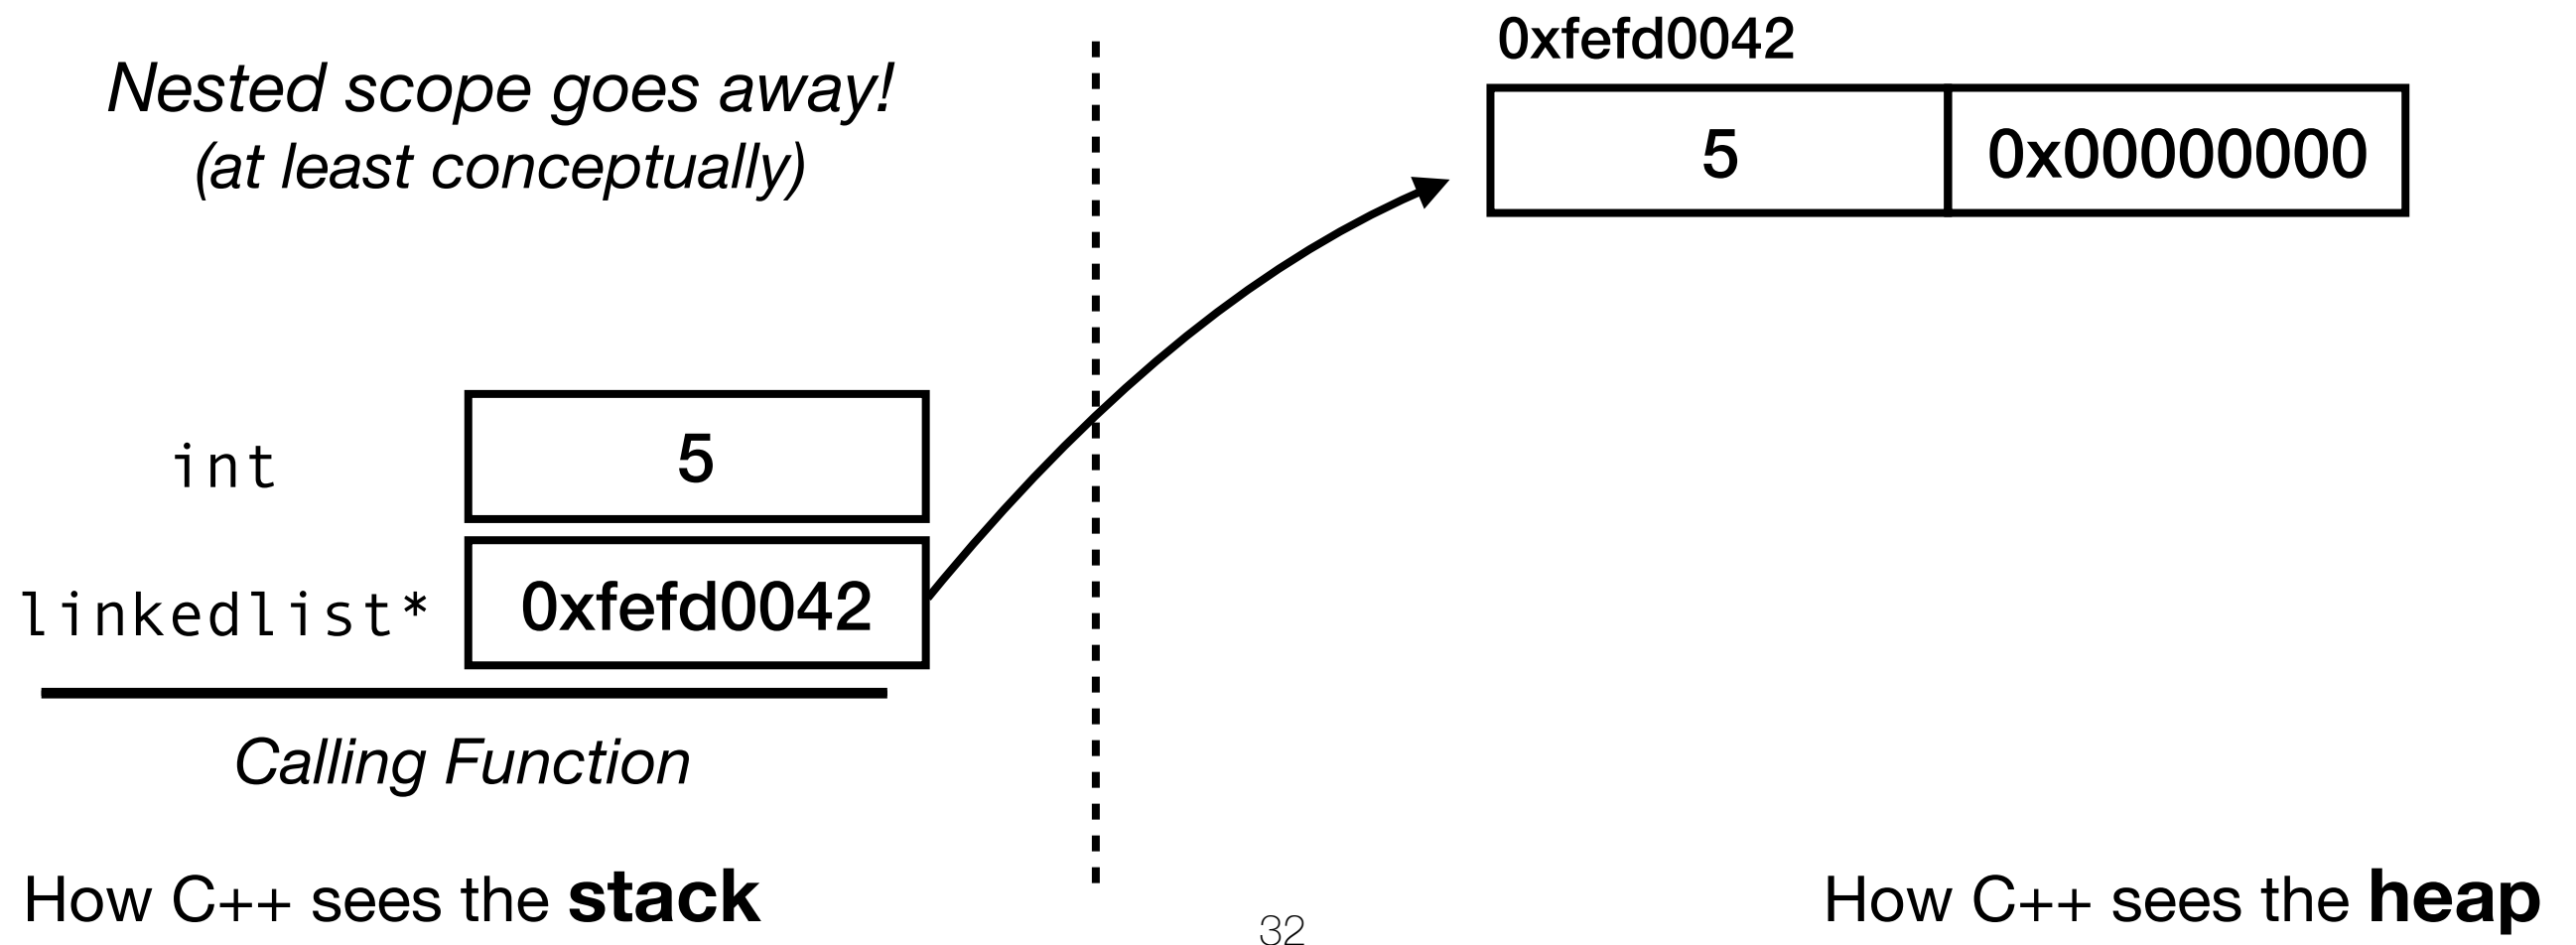
```



```

linkedlist* node = 0; //root
→ int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```

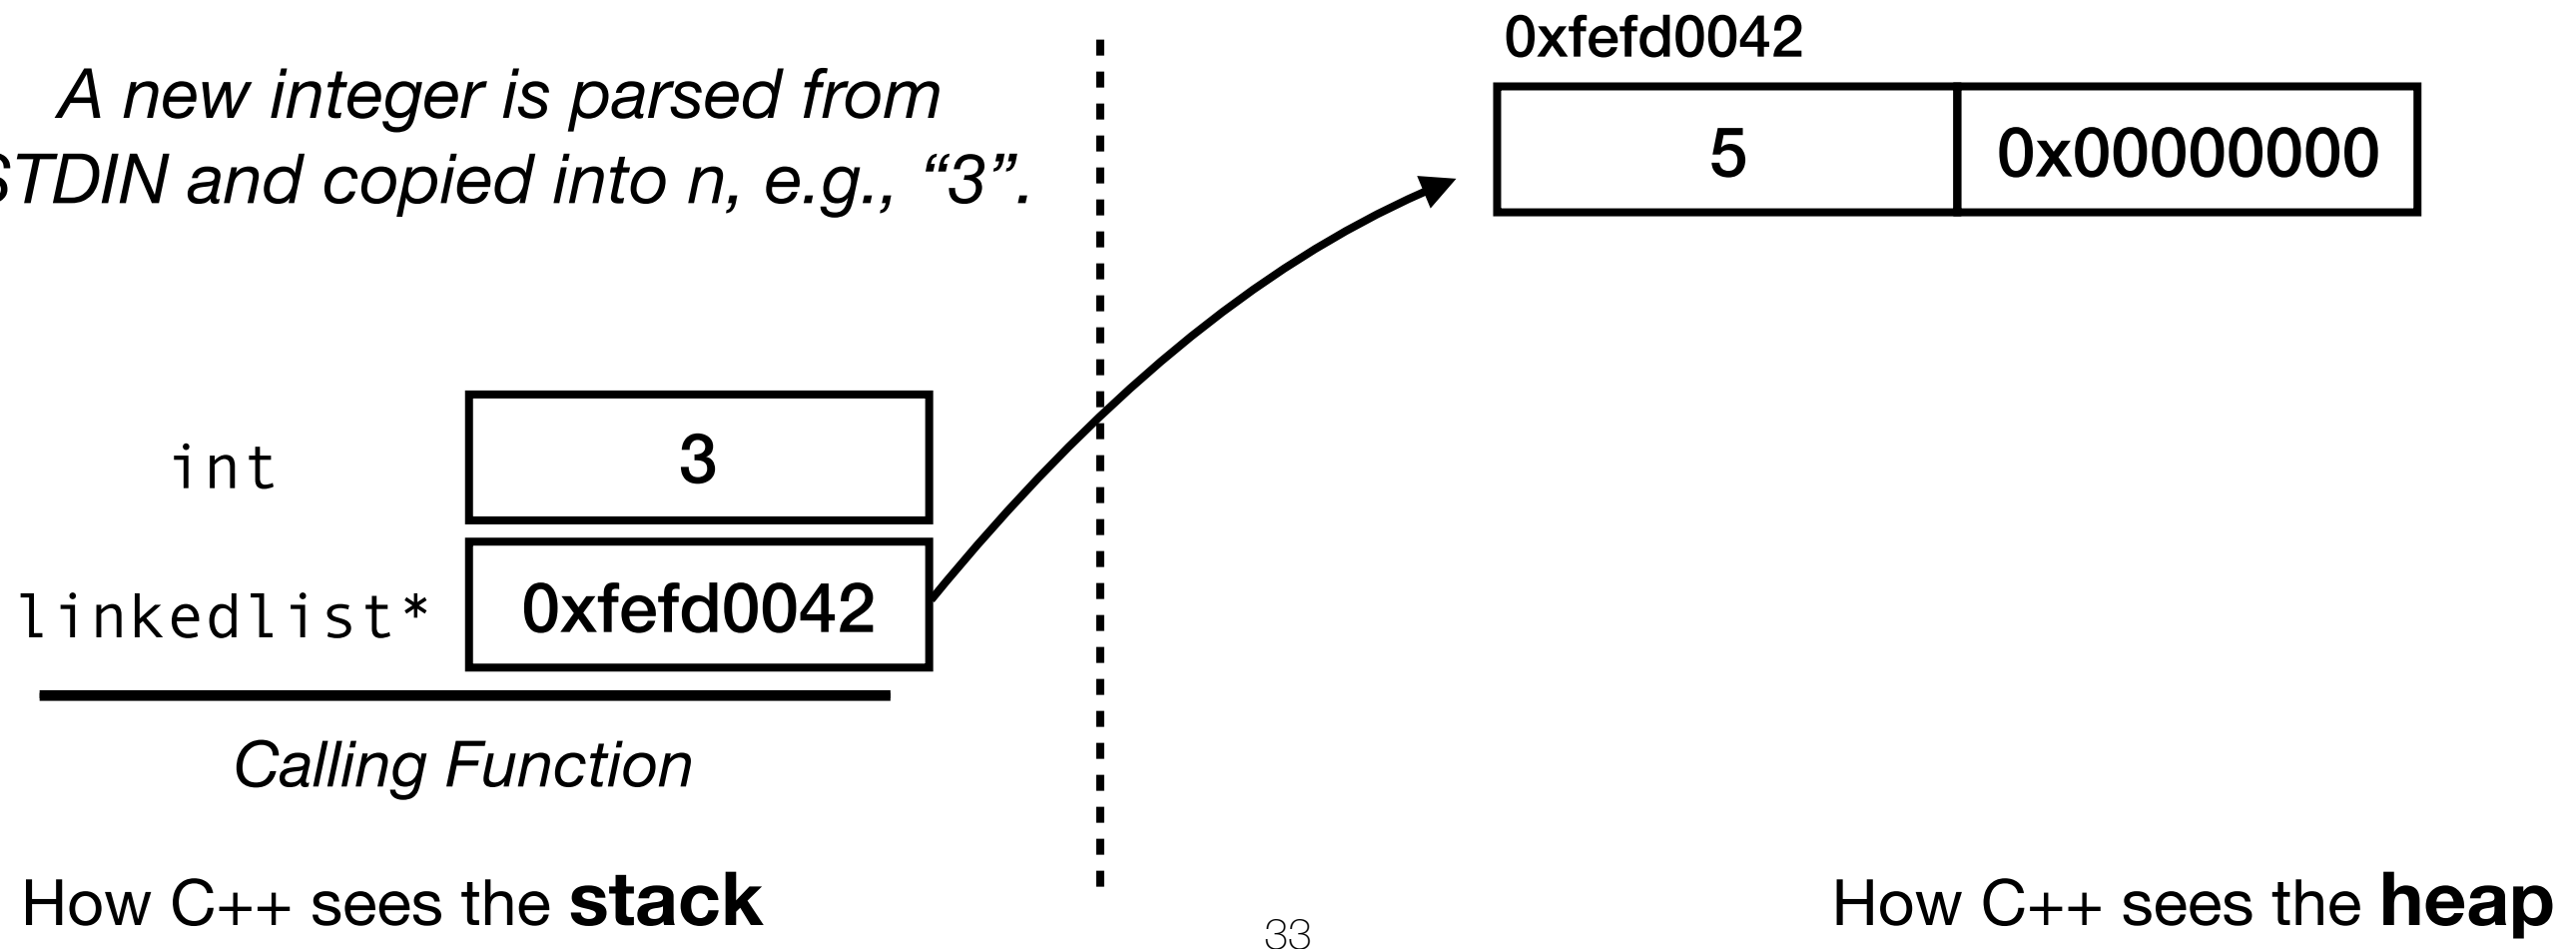



```

    linkedlist* node = 0; //root
    int n;
    while (std::cin >> n)
    {
        linkedlist* next = node;
        node = new linkedlist();
        node->value = n;
        node->next = next;
    }

```



*A new integer is parsed from
STDIN and copied into n, e.g., "3".*



```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}



```



```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}


```



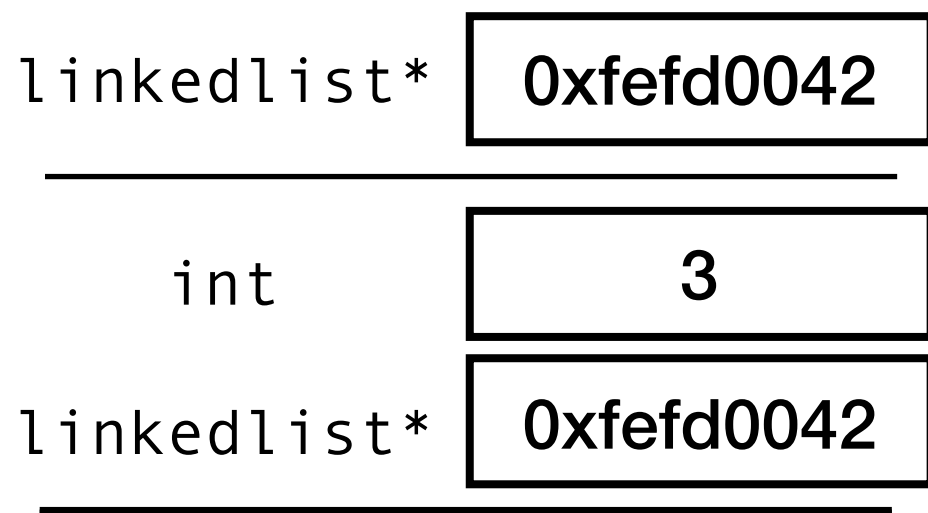
```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```




*assignment
(copies root pointer value)*



Calling Function

How C++ sees the **stack**

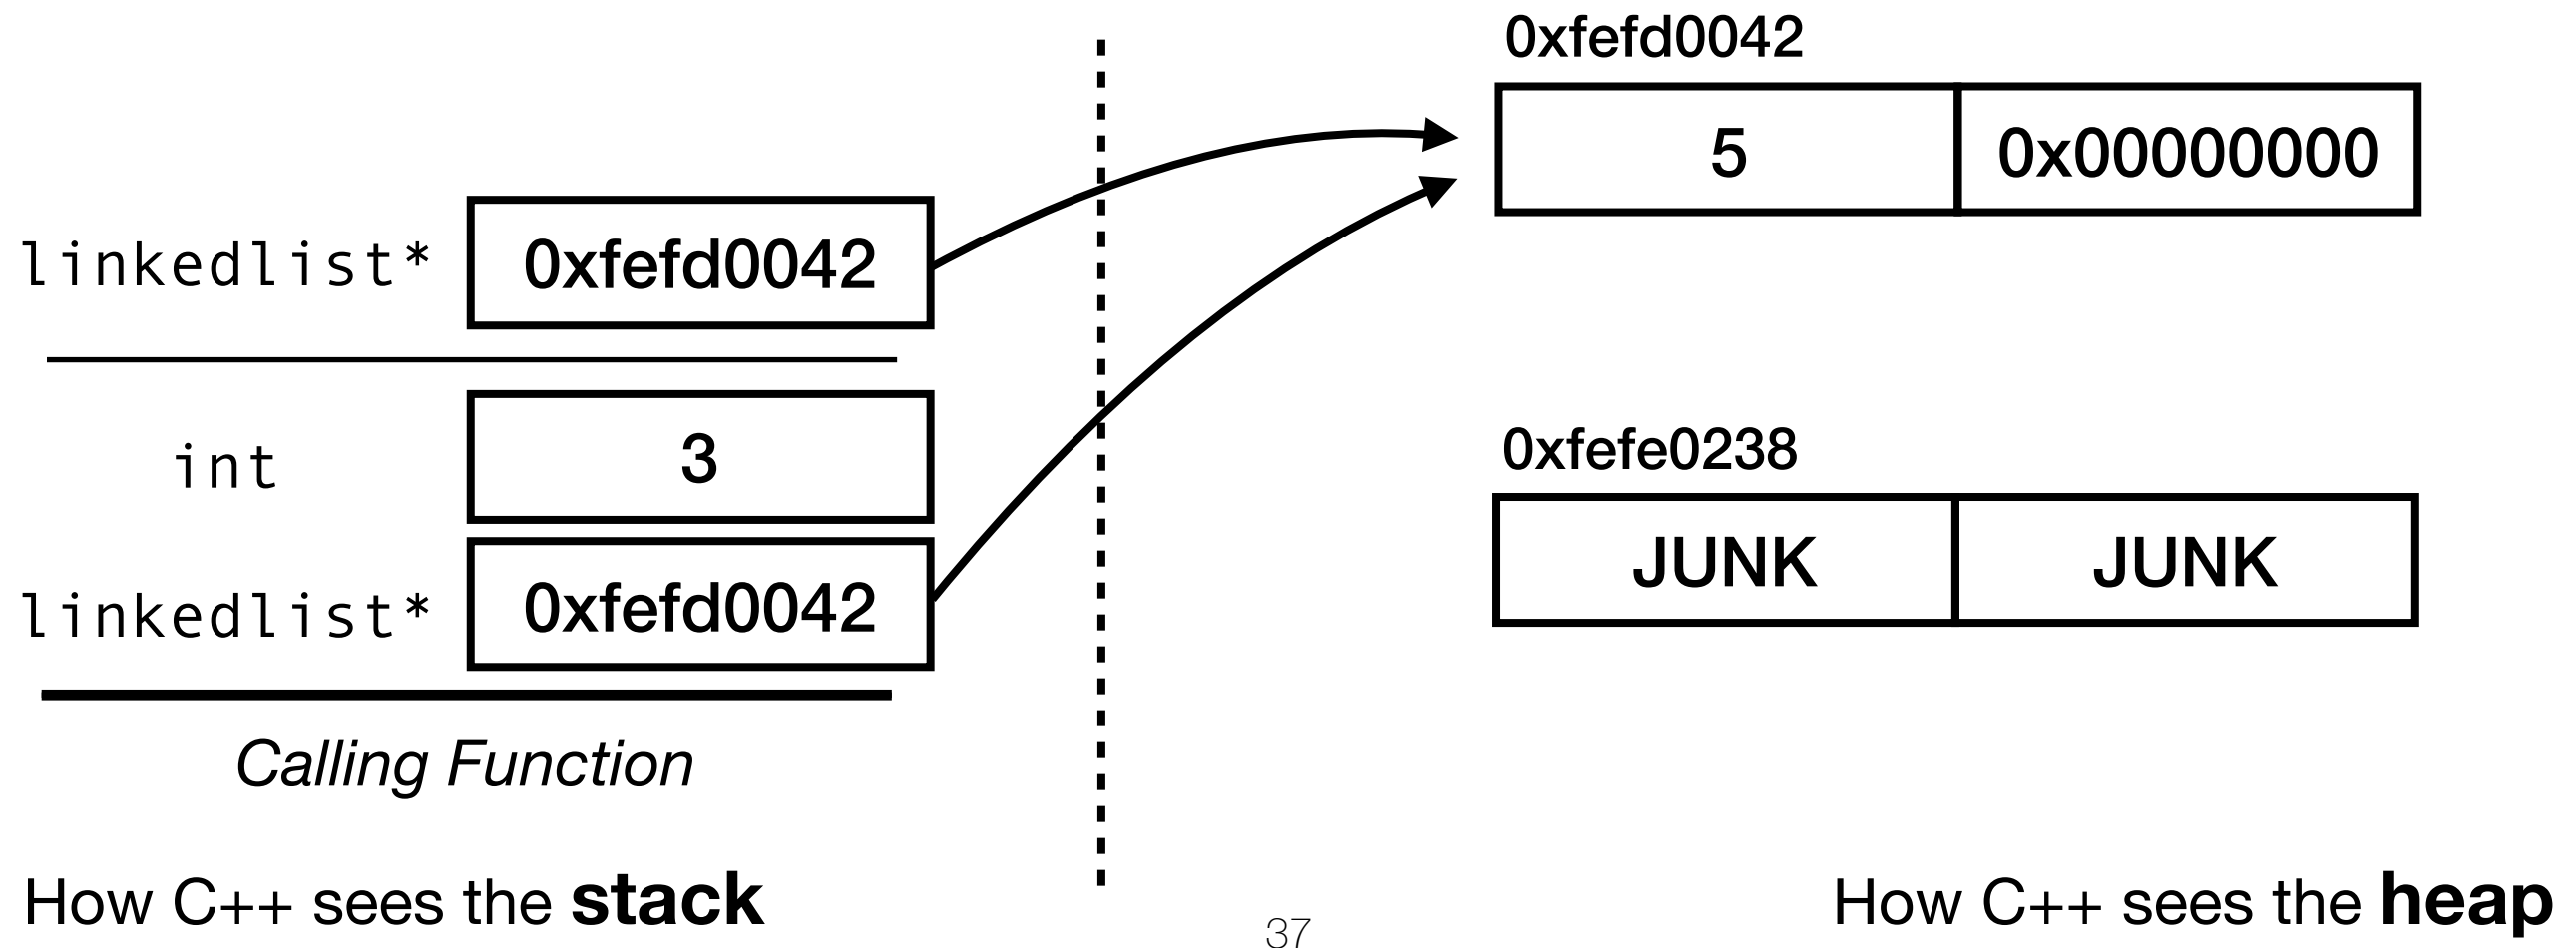



How C++ sees the **heap**

```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}


```



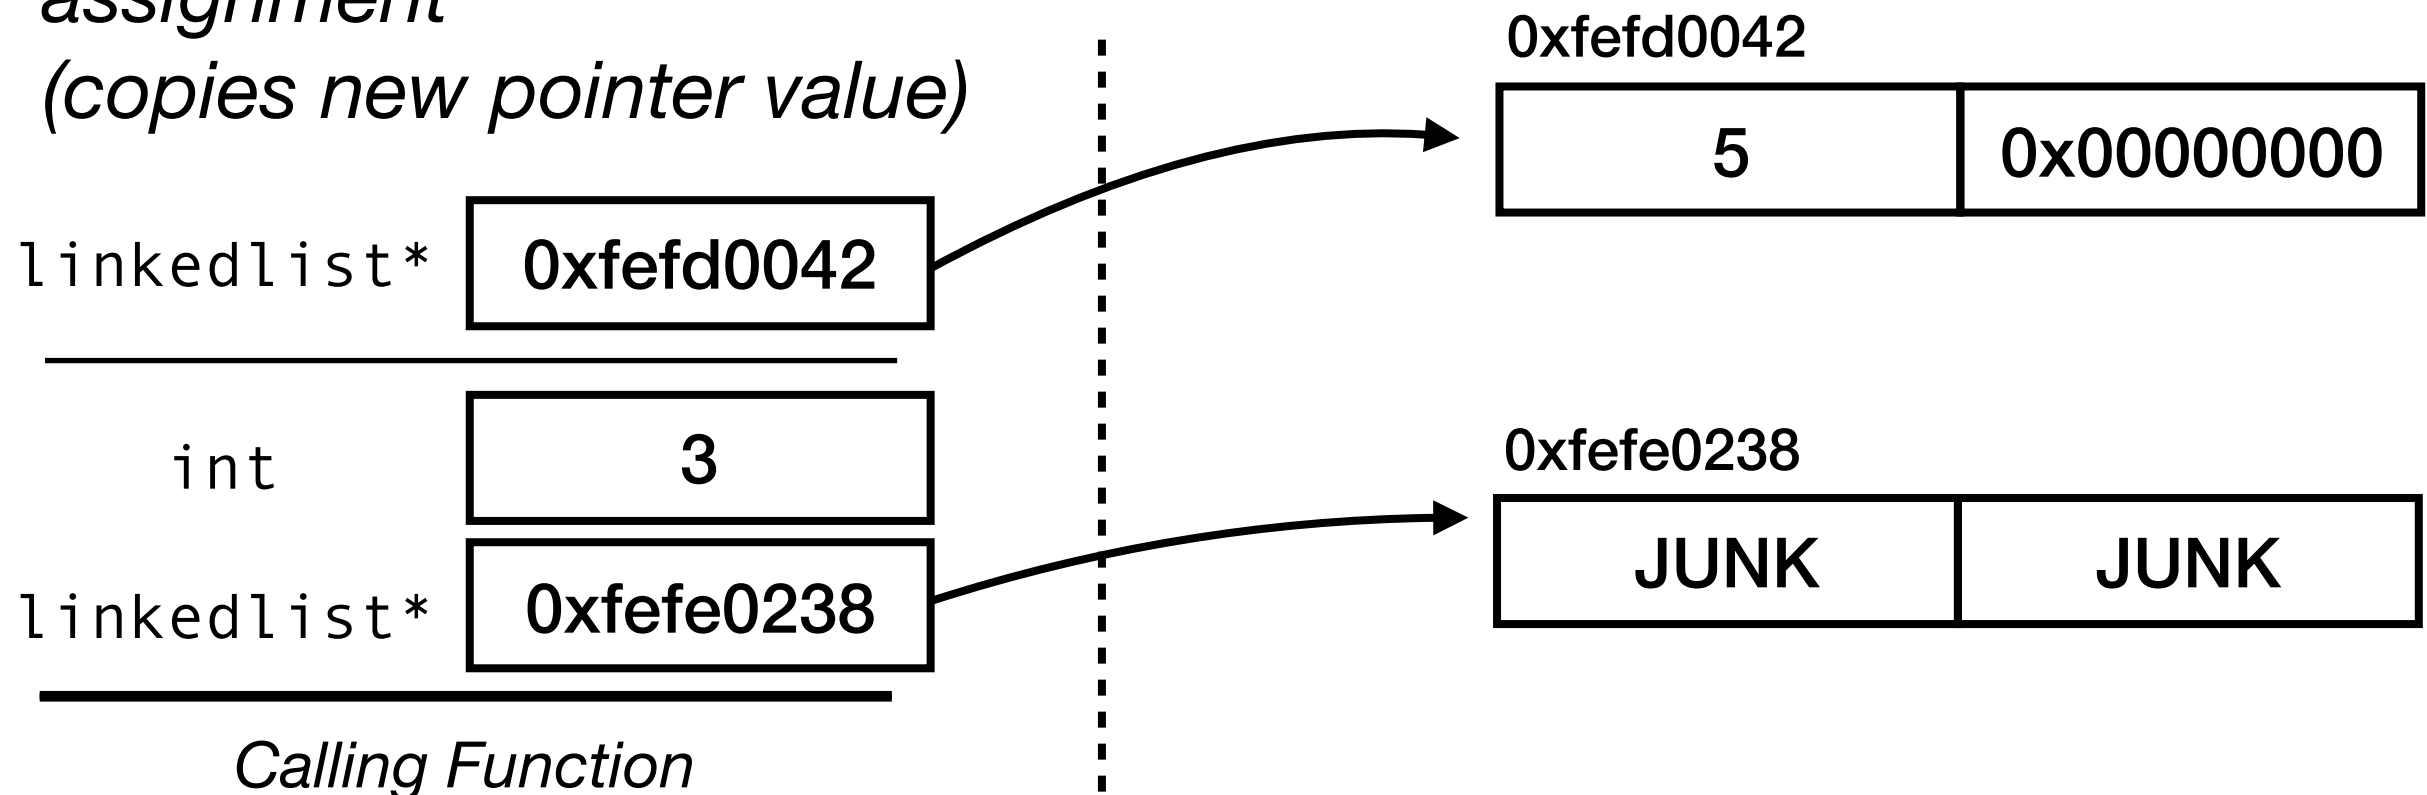
```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```



*assignment
(copies new pointer value)*



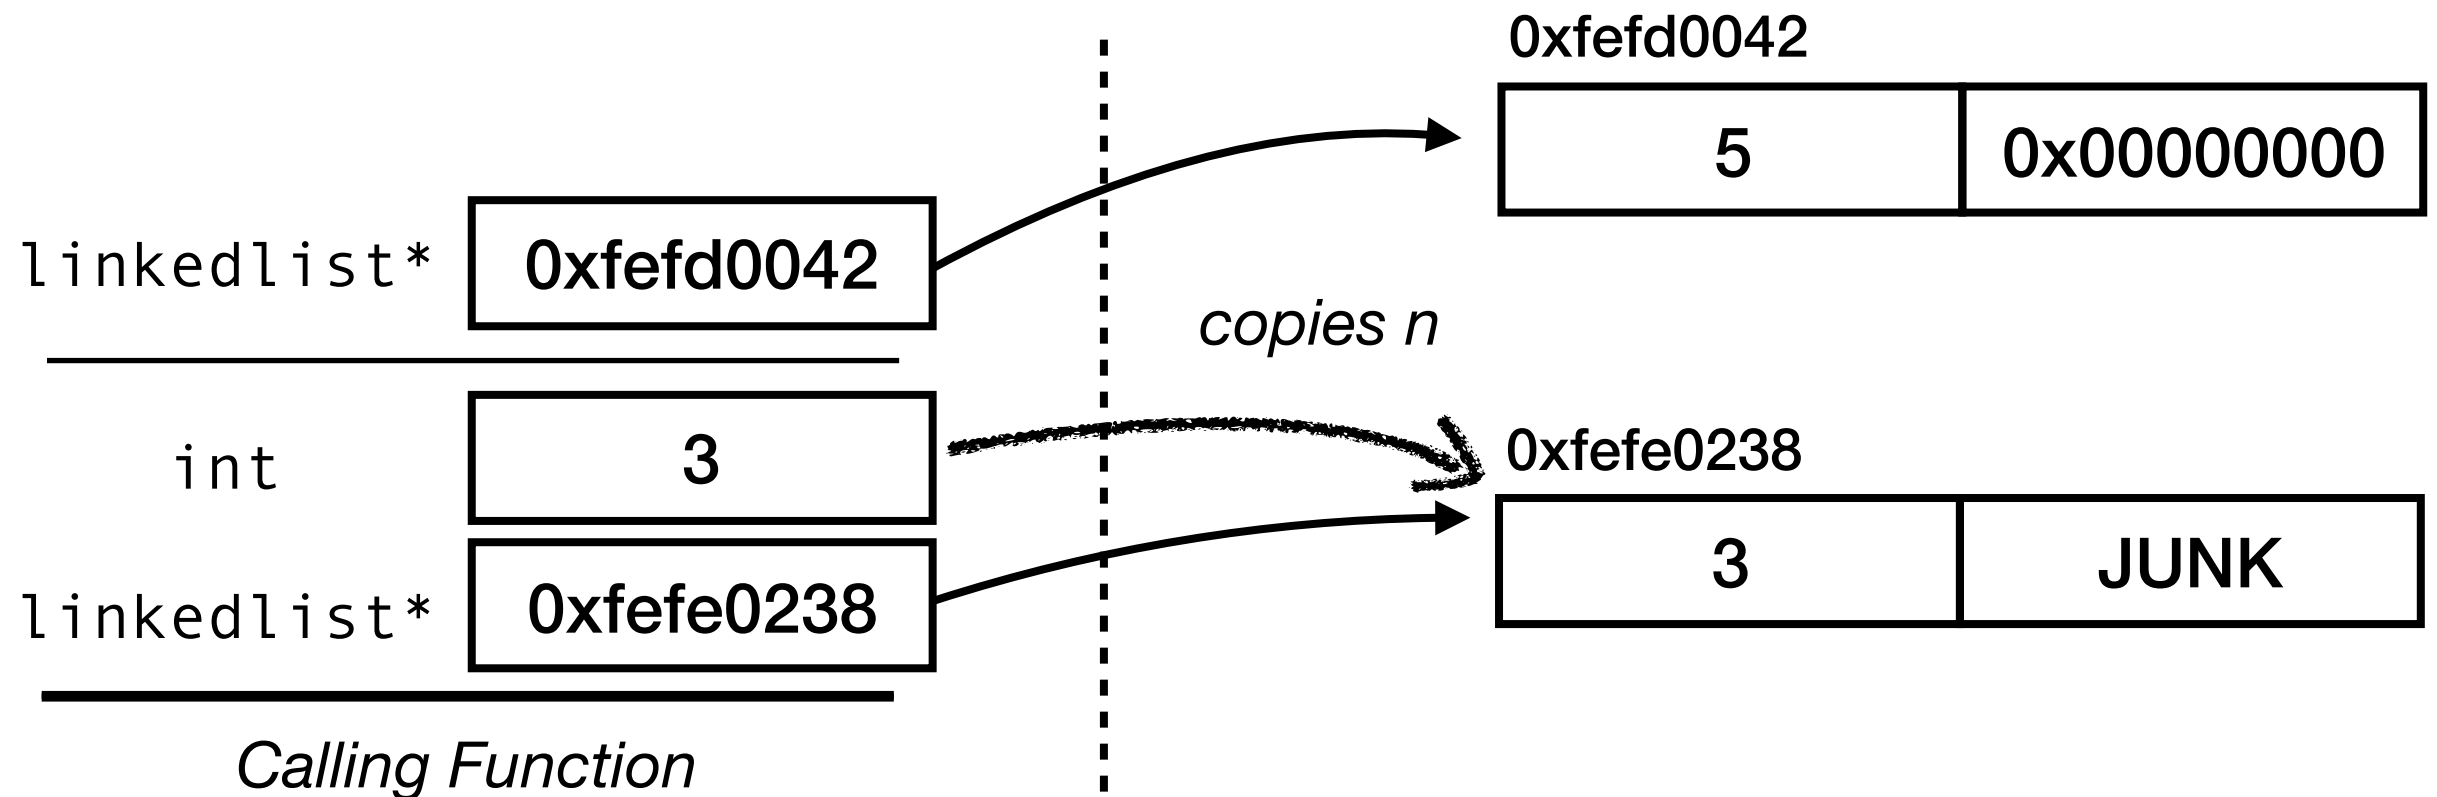

How C++ sees the **stack**

How C++ sees the **heap**

```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```




How C++ sees the **stack**

How C++ sees the **heap**

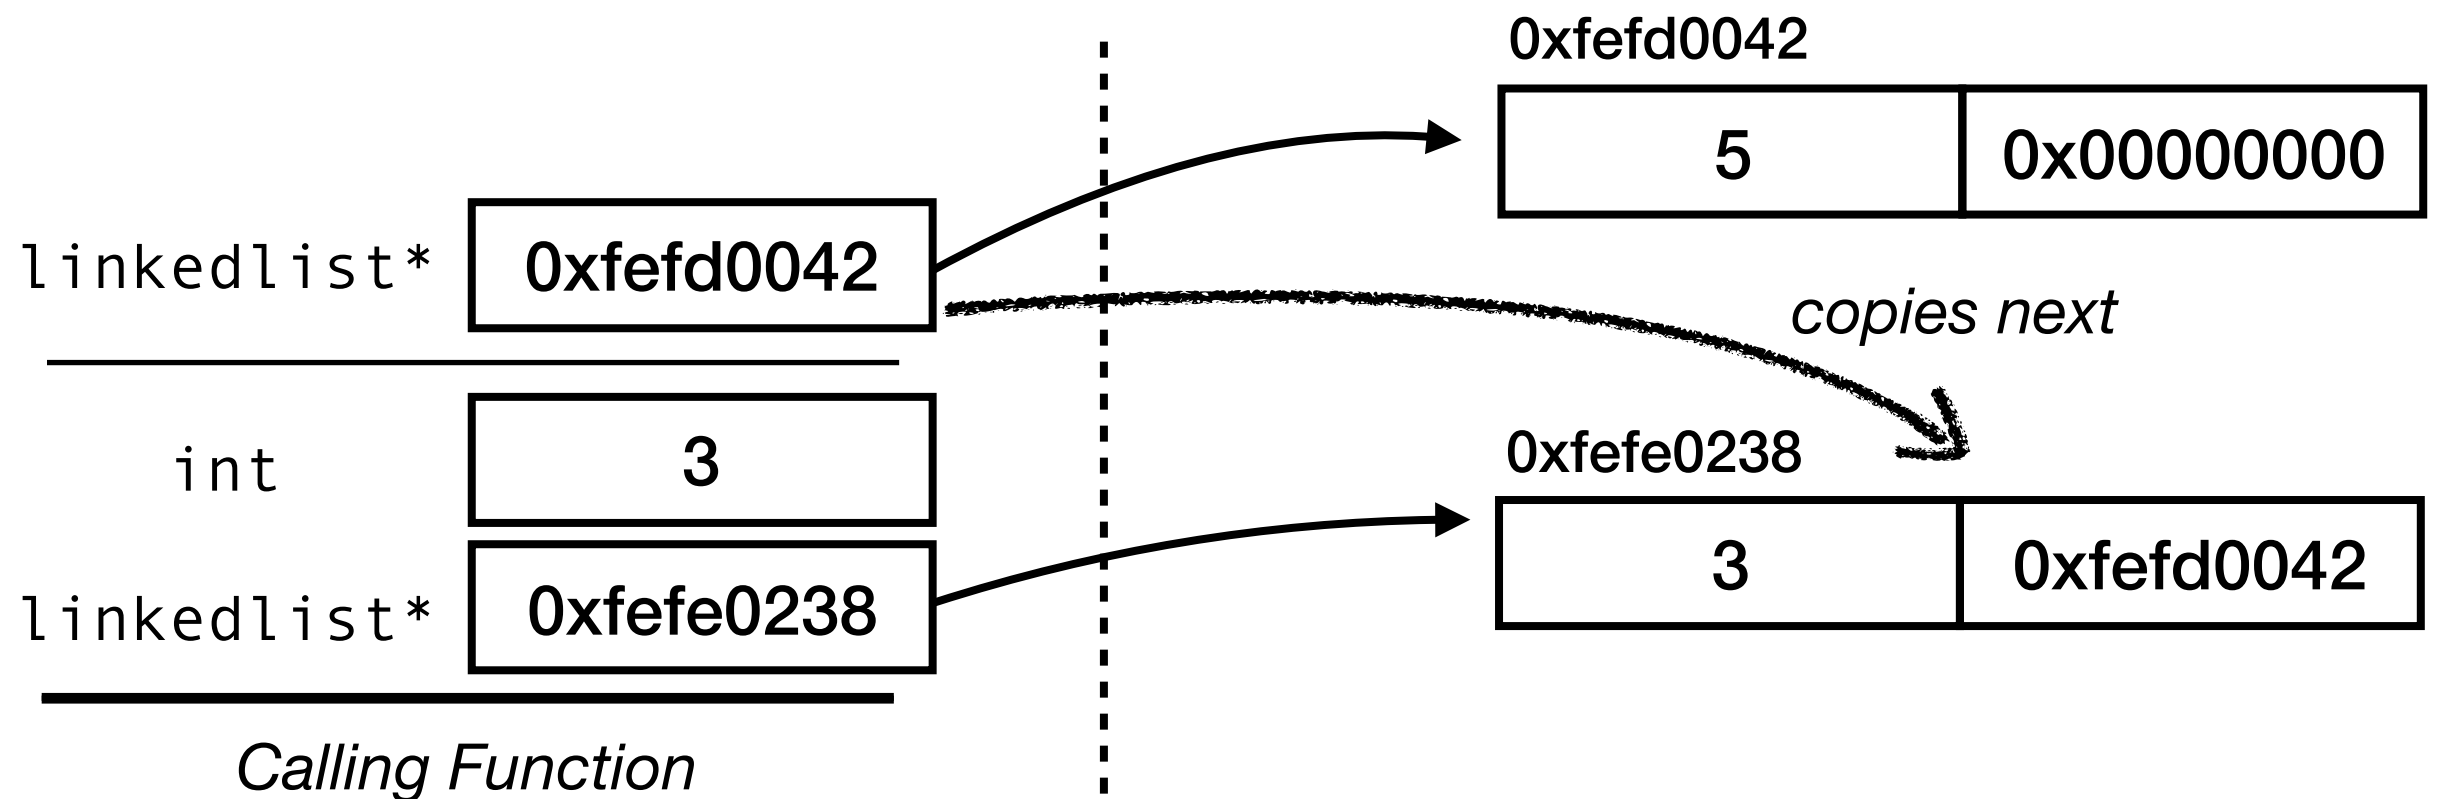
```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```



}



How C++ sees the **stack**


How C++ sees the **heap**


```

linkedlist* node = 0; //root
int n;
→ while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}


```

*Nested scope goes away!
(at least conceptually)*



Calling Function

How C++ sees the **stack**




How C++ sees the **heap**


```

linkedlist* node = 0; //root
int n;
while (std::cin >> n)
{
    linkedlist* next = node;
    node = new linkedlist();
    node->value = n;
    node->next = next;
}

```




(std::cin >> n) reads the EOF (ascii code 0) character and returns false without modifying variable n



Calling Function

How C++ sees the **stack**



How C++ sees the **heap**

C++ semantics: taking pointers of stack values

```
char* badalloc()
{
    char bytes[4096] = {0};
    return &bytes[0];
}

int main()
{
    char* arr = badalloc();
    arr[0] = 'h';
    arr[1] = 'i';
    std::cout << arr << std::endl;
    return 0;
}
```

C++ semantics: Try an example

```
char* badalloc()  
{  
    char bytes[4096] = {0};  
    return &bytes[0];  
}
```

What could go wrong when allocating memory this way?

```
int main()  
{  
    char* arr = badalloc();  
    arr[0] = 'h';  
    arr[1] = 'i';  
    std::cout << arr << std::endl;  
    return 0;  
}
```

C++ semantics: taking pointers of stack values

```
$ clang++ -o bin badalloc.cpp
badalloc.cpp:8:13: warning: address
of stack memory associated with local
variable 'bytes' returned [-Wreturn-
stack-address]
```

```
    return &bytes[0];
           ^~~~~~
```

```
1 warning generated.
```

```
$ ./bin
```

```
hi
```

```
$
```

C++ semantics: taking pointers of stack values

```
char* passthrough(char* ptr)
{
    return ptr;
}
```

```
char* badalloc()
{
    char bytes[4096] {0};
    return passthrough(&bytes[0]);
}
```

```
int main()
{
    ...
}
```

C++ semantics: taking pointers of stack values

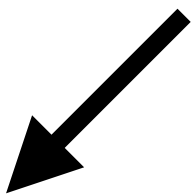
```
$ clang++ -o bin badalloc.cpp  
$ ./bin  
hi  
$
```

The compiler wont always
catch this problem for us!

C++ semantics: taking pointers of stack values

```
char* badalloc()  
{  
    char bytes[8] = {0};  
    return passthrough(&bytes[0]);  
}
```

Now we can try making
the buffer small!



```
int main()  
{  
    char* arr = badalloc();  
    arr[0] = 'h';  
    arr[1] = 'i';  
    std::cout << arr << std::endl;  
    return 0;  
}
```


C++ semantics: taking pointers of stack values

```
$ clang++ -o bin badalloc.cpp  
$ ./bin  
\300I\211\350\376^?  
$
```

Now the call to `std::cout` itself tramples on this stack space and overwrites these bytes with values that are, to us, junk!

How Objects Work

C++ dynamic dispatch: Try an example!

```
class B
{
    virtual int f() { return 1; }
};
class A : public B
{
    virtual int f() { return 2; }
};
```

```
B* a = new A(); // Get a pointer to an A obj
std::cout << a->f() << std::endl;
```

// Which value is printed out?

C++ dynamic dispatch: Try an example!

```
class B
{
    virtual int f() { return 1; }
};
class A : public B
{
    virtual int f() { return 2; }
};
```

```
B* a = new A(); // Get a pointer to an A obj
std::cout << a->f() << std::endl;
```

// Which value is printed out? ANSWER: 2

Function pointers

```
int add1(int x) { return x+1; }
```

In stored-program machines, all code sits somewhere in memory.

In C/C++, you can obtain pointers to functions at run-time, and invoke them! The pointer for `add1` can be obtained with:

`&add1`

```
int add1(int x) { return x+1; }
```

```
int main()  
{  
    int (*f)(int) = &add1;  
  
    // ...  
  
    int four = (*f)(3);  
}
```

**A function pointer, cmp,
passed to sort as an argument.**


```
int sort(int* x, int len, bool (*cmp)(int,int))
{
    // ...

    // ...
    if ((*cmp)(*x,*y))
    {
        swap(*x,*y);
        // ...
    }

    // ...
}
```

**The function pointer, cmp,
dereferenced and invoked.**



```
{  
    // ...  
    sort(buff, length, &lessthan);  
    // ...  
}
```



**A pointer to function `lessthan`
is passed into `sort`.**

C++: Try an example!


Talk to your neighbors. Can you think of another way to parameterize a `sort` method over the comparison predicate to be used?

 **A function pointer, cmp, type int x int -> bool,
is a template parameter to sort.**

```
template <bool (*cmp)(int,int)>
int sort(int* x, int len)
{
    // ...
    if ((*cmp)(*x,*y))
    {
        swap(*x,*y);
        // ...
    }
}
```

**Templated function sort is
invoked with a template
parameter like so: sort<...>(...)**

```
int main()
{
    // ...
    sort<&less<int>>(buff, length);
}
```



C++ dynamic dispatch: class polymorphism

```
class Cmp
{
    virtual bool cmp(int x, int y) = 0;
};
class LessThan : public Cmp
{
    virtual bool cmp(int x, int y)
    { return x < y; }
};
class GreaterThan : public Cmp
{
    virtual bool cmp(int x, int y)
    { return x > y; }
};
```

An instance of type Cmp, cmp, has overloaded method cmp.

```
int sort(int* x, int len, const Cmp& cmp)
{
    // ...
    if (cmp->cmp(*x, *y))
    {
        swap(*x, *y);
        // ...
    }
}
```


**Pass in object less than
by reference to polymorphic
type Cmp supporting the
Cmp::cmp(int, int) member.**

```
int main()
{
    // ...
    LessThan lessThan;
    sort(buff, length, lessThan);
}
```

Virtual Tables (vtables)

Virtual Tables (vtables)

**A table of virtual methods
with a function pointer for each**



```

class Animal
{
    virtual const char* name() = 0;
    virtual int weight() const = 0;
    virtual void eat(Animal* prey)
    {
        if (this->weight()
            < 2 * prey->weight())
            return;
        delete prey;
        std::cout << prey->name()
                  << " was eaten!\n";
    }
};

```



```
class Mouse : public Animal
{
    int grams;

    Mouse(int grams)
        : grams(grams) {}

    virtual const char* name()
    {
        return "Mouse";
    }

    virtual int weight() const
    {
        return this->grams;
    }
};
```

```
class Cat : public Animal
{
    Cat() {}

    virtual const char* name()
    {
        return "Cat";
    }

    virtual int weight() const
    {
        return 4260;
    }
};
```

```

class Giraffe : public Animal
{
    virtual const char* name()
    {
        return "Giraffe";
    }
    virtual int weight() const
    {
        return 1570000;
    }
    virtual void eat(Animal* prey)
    {
        std::cout << this->name()
                    << " wont eat that.\n";
    }
};

```

```
// vtable struct for Animal subclasses
struct AnimalVTable
{
    const char* (*name)(void*);
    int (*weight)(const void*);
    void (*eat)(void*,void*);

    AnimalVTable(const char* (*name)(void*),
                  int (*weight)(const void*),
                  void (*eat)(void*,void*))
        : name(name), weight(weight), eat(eat)
    {}
};

// Allocate a vtable for each concrete Animal
AnimalVTable mouse_vtable(&nameMouse,
                           &weightMouse,
                           &eatAnimal);
```

```
// Class Mouse compiled to a struct
```

```
struct Mouse
```

```
{
```

```
    AnimalVTable* vptr;
```

```
    int grams;
```

```
};
```

```
// An allocator/constructor for Mouse
```

```
Mouse* newMouse(int grams)
```

```
{
```

```
    Mouse* m = (Mouse*)malloc(sizeof(Mouse));
```

```
    m->vptr = &mouse_vtable;
```

```
    m->grams = grams;
```

```
    return m;
```

```
}
```

```
// A name method for Mouse instances
const char* nameMouse(void* _ths)
{
    return "Mouse";
}
```

```
// A weight method for Mouse instances
int weightMouse(const void* _ths)
{
    const Mouse* ths = (const Mouse*)_ths;
    return ths->grams;
}
```

```
// Looks up the vtable for an object
VTable* vtable(void* obj)
{
    return (VTable*)((void**) obj)[0];
}

{
    // To call a member function f:
    // e.g., obj->f(arg0, arg1, ...);

    vtable(obj)->f(obj, arg0, arg1, ...);
}
```

```

// Looks up the vtable for an Animal object
AnimalVTable* vtable(void* obj)
{
    return (AnimalVTable*)((void**) obj)[0];
}

// A default eat method for Animals
void eatAnimal(void* ths, void* prey)
{
    if (vtable(ths)->weight(ths)
        < 2 * vtable(pre)->weight(pre))
        return;
    delete prey; // vtable(pre)->~Animal...
    std::cout << vtable(pre)->name(pre)
               << " was eaten!\n";
}

```


Try an example:

How do you define the constructor
for Giraffe?

```

// Class Giraffe compiled to a struct
struct Giraffe
{
    AnimalVTable* vptr;
    // No data members
};

AnimalVTable giraffe_vtable(&nameGiraffe,
                             &weightGiraffe,
                             &eatGiraffe);

// An allocator/constructor for Giraffe
Giraffe* newGiraffe()
{
    Giraffe* g = new Giraffe();
    g->vptr = giraffe_vtable;
    return g;
}

```

Try an example:

How do you define the virtual member functions for Giraffe?

```

const char* nameGiraffe(void* _ths)
{
    return "Giraffe";
}

int weightGiraffe(const void* _ths)
{
    return 1570000;
}

void eatGiraffe(void* _ths)
{
    Giraffe* ths = (Giraffe*)_ths;
    std::cout << vtable(ths)->name(ths)
               << " wont eat that.\n";
}

```